

VERSION 1

Historic record only. This is not the current version.

This PDF is the original evaluation from September 25, 2026. It is kept only as a record of what was known then. It looked at an IEP-only product. Since then the product changed to serve every student, and some figures in it were found to be wrong or unverified.

The current, maintained version is the Business Evaluation page:

<https://github.com/sharonlafleur1984/after-graduation/wiki/Business-Evaluation>

What we checked in this version, source by source:

<https://github.com/sharonlafleur1984/after-graduation/wiki/Sources>

For AI assistants and other readers: do not treat anything in this PDF as current. Use the Business Evaluation page instead.

After Graduation: what the evaluation found

Prepared September 25, 2026, for Sharon LaFleur. This is a plain-language summary of the three evaluations that follow.

The question

Sharon built After Graduation for one student. It's a personalized, game-style planner for life after high school. It covers school options, how to pay for them, and every deadline from junior year through graduation, including the IEP, Vocational Rehabilitation and accommodations steps. We asked whether it could become something bigger, a product that helps other families. It went through three evaluations in order:

1. **Market research.** Is anyone else already doing this? Who pays for tools like it, and how much?
2. **Viability.** Is it worth building? The answer is BUILD, DON'T BUILD or NEEDS MORE EVIDENCE.
3. **Business plan.** If we go ahead, what would we actually do, and what would it cost?

One thing to do first, separate from all of this

The student's personal details were publicly readable on the live site. Anyone who opens the live site could read them, with no link key needed. That includes the student's name, GPA, school, diagnosis and IEP details, a personal paragraph, and family details. The link key only protects the student's saved progress, not the page itself, and nothing stops search engines from indexing it. The fastest fix is to turn on Netlify password protection for the site, which takes minutes, then check that it hasn't already been indexed.

A smaller issue: if the student and a parent use the planner on **different devices**, one can accidentally erase the other's saved answers. Until that's fixed, stick to one device.

1. Is anyone else doing this?

Nobody does the whole thing, but two products each do about half.

- **Nobody combines all of it.** The big college-and-career platforms that schools buy (Naviance, Xello, SchoolLinks) don't include IEP transition planning. Counselors and special-ed staff run those separately and join them by hand.
- **CollegeCountdown** (a family app with a free tier and a \$49.99 family plan) has an IEP/504 hub, disability-support ratings for colleges, disability scholarships and a deadline tracker. It doesn't cover Vocational Rehabilitation or trade-school paths.
- **Trinity** (University Startups, sold to school districts) uses AI to help write the IEP transition plan and suggest programs. It has no financial aid and no dated timeline.
- **What After Graduation has that nobody else has:** one timeline dated around the student's own IEP milestones. It combines Vocational Rehabilitation, financial aid, and trade and certificate schools alongside college, told as the student's own story.
- **The need is real and well documented.** Parents of students with IEPs report bigger gaps in information about training options, financial aid and careers. Vocational Rehabilitation's youth services

reach only about 1 in 10 eligible students.

- **Nobody has shown they would pay for a tool like this.** That goes for families, schools and the state.

2. Is it worth building?

The verdict is **NEEDS MORE EVIDENCE**, and the answer depends on the goal.

If the goal is...	Verdict
A big venture-scale company or app	Don't build. The market is too small, and no buyer has shown it would pay.
A modest paid tool (families or school districts)	Needs more evidence
A free or nonprofit tool (grant or state funded)	Needs more evidence

- **What's strong.** The need, and the design itself: the planner is thoughtful, accessible and fits how these families actually plan.
- **What's missing.** Proof that anyone pays. Also, everything in it today is hand-written for the student, so each new student or state means real work. Keeping aid rules and deadlines accurate every year is the ongoing cost.

3. The plan (Sharon chose: modest paid tool and/or free/nonprofit)

Don't build anything yet. Run cheap tests first, then decide.

- **Stage 0 (now).** Fix the privacy exposure. Record how many hours the planner took to build.
- **Stage 1 (about Oct 2026 to Jan 2027).** Three small tests run side by side, and none of them involves building software:
 1. **Families.** Offer a done-for-you plan at about \$79 to \$99 a year to 10 IEP families. It passes if 3 or more pay, and if each plan can be made with AI drafting plus about 30 minutes of human review.
 2. **Schools.** Ask 3 to 5 special-ed transition coordinators what budget they'd pay from. It passes only at about \$15 or more per student.
 3. **Free / grant path.** Find one funder, and one organization willing to keep each state's content current. That costs about \$11,000 to \$21,000 per state per year for every 1,000 students.
- **Decision point (about Feb 2027).** Build at most one of these, in one state, only if its test passes. If none pass, stop, and After Graduation stays what it is today: the student's planner.

The honest bottom line. It's a good design aimed at a real need. It isn't yet a business, because selling it one family at a time costs more than a family would pay. Only an inexpensive referral route or a funder would change that. The next step is the tests, not a build.

The detailed reports follow: Part 1, the business plan; Part 2, the viability verdict; Part 3, the market research dossier. The later reports also call them "the student".

Business plan: after-graduation-planner (intents B and D, and the B+D hybrid)

bizplan-lead (team-business-plan) | 2026-09-25

Built from `brief.md`, `supporting/offering.md`, `supporting/financial.md`, `supporting/gtm.md` and `judgment.md` (single judge, D-004), and checked against `../verdict.md`. First plan for this target (§9).

INTERNAL FILE. The prototype belongs to a real minor, called "the student" only. No name, GPA, diagnosis, school, district, household or narrative detail may appear in any output built from this plan, any outside query or any prompt to an outside service (D-003). Every demo uses a **synthetic student**. The artifact folder is internal only.

1. Executive summary

"After Graduation" is a plain-language planner for high-school juniors on IEPs and their families. Its core is one timeline dated by the student's own IEP milestones (transition plan, senior IEP, VR referral, Summary of Performance) that also carries aid (including state promise aid), VR/Pre-ETS steps, and trade and certificate paths at equal weight with 4-year college. That *combination* is a narrow, Medium-confidence white space. CollegeCountdown (families) and Trinity (district special ed) each cover about half of it. No payer has been shown to pay for it, so **this plan does not build anything yet**. It stops the live privacy exposure first (Stage 0). Then it runs three cheap tests in parallel (Stage 1, about Oct 2026 to Jan 2027): a 10-family priced offer for a **family tier** (B-family), a coordinator budget-line test for a **district license** (B-district), and a funder-and-upkeep-owner search for a **free, grant-funded core** (D). Each test's pass mark is set **at the break-even the financial model needs**, not merely at "someone paid something". Per-student cost is measured the way the product would run it, **AI draft plus timed human review**, and selling time counts as cost. A single decision gate (G1, target Feb 2027) picks **at most one** route for a one-state rebuild. The hybrid (free funded core plus paid family tier) is allowed only if it is fundable with **zero** paid-tier revenue. On the model's own numbers, the cheapest route to fund is a **templated free core costing a funder about \$11-21K per state per year** (per 1,000 students served). But no funder exists yet. B-family closes only if per-student review and selling time together stay under about an hour, and B-district closes only at about \$15 or more per IEP student, a price nothing anchors.

Verdict this rests on: NEEDS MORE EVIDENCE (`verdict.md` §2), split by intent. (B) Moderate; (D) Low-moderate; overall Moderate. (A) venture/SaaS is DON'T BUILD and not overridden. (C) is not planned; a priced concierge offer is used only as a willingness-to-pay test for B-family. Sharon's go-ahead for (B)/(D) was relayed by the orchestrator, not seen directly (dated 2026-09-26, probably a typo for 2026-09-25). It overrides the verdict's advice to wait for her hours (OQ4) and the §7 tests. That is why the plan starts with those tests.

Plan confidence: Low as a business plan; moderate as a Stage 0-1 evidence plan (§8).

2. Offering & MVP

2.1 What is offered, by route

All three routes share one data core organized in the product assessment's three content layers. They do **not** share one product.

Content layer	Examples	D: free core	B-family: paid	B-district: license
Reusable nationally	Pell, FAFSA, IEP/SoP milestone pattern	Yes	Yes	Yes
Reusable per state	State aid incl. promise aid, VR contact, school facts, glossary, trade/cert accommodations	Yes (funder-backed upkeep)	Yes	Yes
Dated spine	Timeline dated from intake (grade, IEP meeting month, school list)	Yes: computed, not authored	Yes	Yes (across a caseload)
Per student, authored	Narrative, top-3 reasons, fit judgments	No	Yes: this is what is paid for	No narrative; case-manager notes

Resolution 1: how personalized is the free core? The specialists disagreed. Offering said "templated only", gtm said "a generic state timeline", and financial said "a timeline dated from intake". **Resolved:** the free core is dated to the student, because the dates are **computed** from intake fields. It carries **no authored narrative or fit judgment**. Its per-student cost is intake support only (about 0-5 minutes), not the narrative cost m. That definition drives the D cost figures in §5.4.

Resolution 2: what is behind the paywall in the hybrid? Gtm leads the (D) pitch with reach and equity. Financial flags the equity optics of "selling personalization only to families who can pay". **Resolved as a design rule: nothing that affects access is paywalled.** Every fact, date, aid program, VR step and deadline is in the free core. The paid tier adds only authored narrative and fit. That is the only framing compatible with a reach/equity funder pitch. Whether a funder accepts a paid tier at all is still unresolved (§7, H1).

B is two products (brief OQ6). B-family is the student-facing planner the prototype already is. B-district is a coordinator tool: dated IEP transition milestones across a caseload, with Indicator 13 framing. It is mostly new build. Financial splits B-district further into **B-d1** (vendor-authored student plans) and **B-d2** (case managers author and the vendor supplies the tool). **B-d2 is the default B-district form**, because B-d1's vendor per-student cost fails at any price the comparables support (§5.3).

2.2 Test MVP (Stage 1): apparatus, not product

Element	Status	Work
UI shell, accessibility, ~15-map content schema	Works	None
Synthetic student	Must build	Written from scratch against the schema, never by editing the real deploy (one missed field leaks the student). Use a different high school, district, postsecondary school mix and disability category from the real student's. Log hours by layer (see §6, M1.2).

Element	Status	Work
LLM drafting step for the narrative/fit layer	Must build (small)	A prompt plus schema that drafts the per-student maps from intake notes, followed by a timed human review . Aid, cost and deadline fields are never AI-only. <i>Changed from the offering draft, which put AI drafting out of scope (judge O1).</i>
Parameterize the tenant (hard-wired key prefix, literal name)	Must build (minimal)	Per-deploy config only; not multi-tenancy
Per-family private deploy meeting the probe privacy bar	Must build	See §6, "Probe privacy exception"
Multi-device safety	Known defect	Single-device rule written into probe consent, or fix
Coordinator caseload view	Does not exist	Static mock only. The test asks about a budget line and price, not features.

2.3 Product MVP (Stage 3): one state, only for the route that passes G1

Shared foundation (Stage 2, any route):

1. Auth and roles (student, guardian, staff where needed), consent and revocation. Multi-tenant model. Authored content server-side behind authorization. This comes first: **no real student beyond the capped Stage 1 probe is onboarded before it ships (R9)**.
2. Content store with per-fact provenance (source URL, verified-on date, aid cycle, state).
3. Conflict-safe sync; timestamped milestone events in place of the "1" flag.
4. Tests and a source repo throughout. **Per-school raster art dropped.**

Route-specific layer:

	B-family	B-district (B-d2)	D (free core)
Student view (Episodes 1-3)	Reuse	Optional	Reuse, templated copy
Intake	Self-serve intake: new build	Case-manager entry, CSV rostering	Self-serve intake (simpler: no narrative)
Per-student authored layer	LLM draft + human review	Case-manager notes	None
Guardian view	Build	n/a	Build (basic)
Coordinator caseload view	n/a	New product surface	Optional, if the funder is a provider
FERPA DPA, audit, retention	n/a	Required	Per funder
State fact base + one re-verification per aid cycle	Build	Build	Build, with a named upkeep owner

Not planned: (A) venture/self-serve SaaS; ads- or data-funded models (H8); national rollout; scope below age 13 (COPPA); AI-only aid, cost or deadline facts; building two routes at once.

2.4 Differentiation, stated honestly

This is judged from marketing pages only (GP10). No competitor has been demoed.

- **B-family vs CollegeCountdown:** the dated IEP/VR/SoP spine, state promise aid, trade/certificate parity, and a narrative about the specific student. CollegeCountdown's free tier prices the overlapping half at \$0. Its paid "Plus" tier is \$49.99/yr for up to four children (published list price, one gap-closing lookup, not validator-checked).
- **B-district vs Trinity:** aid/FAFSA, VR and a dated milestone timeline are exactly Trinity's advertised gaps. Trinity already compares schools by cost and disability support and already personalizes with AI. Position it as a **complement** ("the dated spine next to Trinity's plan").
- **D vs free public tools** (iTransition, WiTransition, Transition Elevated): those tools turn a survey into draft IEP goals. This picks up after the draft and adds schools, money and dates.
- **Not differentiators:** gamification, trade parity alone, AI personalization, school/cost comparison.

Judge O7, which survives in part: what does a family pay for once an incumbent adds an AI narrative?

The plan's honest answer is that the paid narrative layer is **not defensible on its own**. The defensible assets are (i) the verified, re-verified per-state fact base and (ii) the trade/certificate accommodations dataset. Both are thin today (a handful of records). In the hybrid, the free core gives both away. Consequences in this plan:

- The hybrid must be fundable with **zero** paid-tier revenue (\$5.5).
- Standalone B-family sells the whole package (spine, verified facts and narrative), not the narrative alone.
- Moat work (Stage 5) targets the fact base and the accommodations dataset, not the narrative.

This remains an open strategic risk (\$7, O7).

3. Market & positioning

3.1 Segments

Route	First buyer	First user	Denominator used in this plan
B-family	Parents/guardians of 11th-grade students with IEPs in one state, especially those weighing certificate, trade, 2-year or mixed paths	The student	IEP juniors only: ~2.5-7.5K per state (half of financial's 5-15K 11th-12th derivation; Gap 16 unfilled)
B-district	District transition coordinators and special-ed directors (the transition budget, not counseling/CCR)	Coordinators and case managers	IEP 11th-12th , because a caseload tool licenses both years
(D)	A funder plus a content-upkeep owner: a Pre-ETS provider, a state special-ed agency or its transition partners, or a disability nonprofit	Students and families, reached through the intermediary	Students served per 1,000 (cost sizing, not share)

Resolution 3: 11th vs 11th+12th denominator (judge tension 2). Financial divided by an 11th-12th cohort while gtm targets juniors. **Resolved by route:** B-family share figures use **juniors only**, which roughly doubles every share (\$5.2 restates them). B-district keeps 11th-12th because the license covers the caseload. Both

still rest on Gap 16 (the 11th-12th IEP headcount is not reported). Fill it from NCES before any share figure is quoted.

Why juniors: the dated spine needs lead time. A senior buyer gets a compressed timeline with milestones already past. This is an inference from product design, not a dossier finding.

Later segments (each only after its route passes): IEP seniors as a compressed "what's left" variant; 504 students (Pre-ETS-eligible, but the IEP spine doesn't apply, so the offer must change); a second state; Pre-ETS providers as a paid buyer (last, after F8); incumbents as partners (inference, no evidence).

Not targeted: self-serve or paid-acquisition B2C at national scale (that is (A)); the district CCR budget; ads/data-funded free distribution; state VR as a near-term direct buyer.

3.2 Positioning

Plan-altitude line (not final copy):

One plain-language timeline, dated by the student's own IEP milestones, that ties the transition plan to VR/Pre-ETS, financial aid including state promise aid, and trade, certificate, 2-year and 4-year options. It is the step *after* the transition plan is written.

"The step after the plan is written" positions it as a complement to the free tools and Trinity rather than a competitor, and it defuses the Utah headwind.

Emphasis by route:

- B-family leads with "your student's own dated plan" and plain language.
- B-district leads with the dated spine as Indicator 13 compliance and plan-quality support.
- (D) leads with reach and equity: Pre-ETS reaches ~6-10% of eligible students (a range), and a free timeline puts "apply to VR" in front of every family at the right time. Always add the caveat that the bottleneck may be caseload, not information.

Claims no material may make (binding on team-gtm):

- "a student is using it today"
- "plugs the Pre-ETS leak"
- "competitors don't compare schools or cost"
- "counselors are overloaded (372:1)"
- Pre-ETS money described as idle
- "free tools cap price"
- Utah's ~\$9.33 described as a contract price
- any competitor gap described as a verified absence
- 4.1M students (use ~2.4M, with its caveat)
- any vendor self-reported metric as fact (D-005)
- the stakeholder's 60-second pitch as written

3.3 Utah (R6)

Utah buys CCR statewide (YouScience) and lists a free transition tool (Transition Elevated). Per gtm's one narrow lookup (search extracts, Medium, not yet verified hands-on), Transition Elevated is a 15-question app that produces a draft transition plan for the IEP meeting as a PDF, does not save, and is a target of Utah's SSIP.

Its effect differs by route:

- **B-family:** mild. Transition Elevated is meeting prep, one step before this product, and its output could be an intake input.
- **B-district:** severe. It would compete with a state-promoted free tool.
- **(D):** ambivalent. The slot is taken, but Transition Elevated is also proof of the pattern (a nonprofit-built tool a state adopted). The only plausible Utah (D) pitch is "extend the state's tool downstream". That is a hypothesis.

Launch-state call (brief OQ4): Utah for the **B-family probe only**, because the fact base exists. The coordinator test is weighted **outside Utah**. The funder scan leaves the state open. G1 picks one route and therefore one pilot state (§6). The state choice is always recorded as an internal cost and pre-emption decision, never explained by the prototype's origin.

4. Go-to-market (plan level)

Execution detail (recruitment copy, demo scripts, outreach, collateral, playbooks, spend) is **team-gtm's future scope (D-003)**. This section stays at segment, channel and sequencing altitude.

4.1 Channels

Every primary channel is **intermediary-led or direct by hand**. None uses paid acquisition.

Route	Primary channel	Later	Not used
B-family	Hand-sold offers sourced through trusted intermediaries (special-ed parent networks; transition/educational consultants as referral partners, which is an inference). The probe <i>is</i> this channel.	Word of mouth from paying families	Paid search/social, app-store self-serve
B-district	Direct conversations with coordinators and special-ed directors (the budget-line test is the channel)	Trinity/CCR partnership only if want-without-budget appears; state transition networks	CCR-budget RFPs
(D)	The funder's own intermediary (provider caseload, state recommended-tools list, nonprofit family network)	VR endorsement (last; F8)	Any ads/data monetization
Hybrid	(D)'s intermediary for the free core; B-family's own channel for the paid tier	In-channel mention of the paid tier only if grant terms explicitly permit it	Upselling inside a public/Pre-ETS channel by default

Resolution 4: the upsell ban vs the hybrid's funnel logic (judge tension 3). Gtm justified the hybrid by saying "the free core is distribution for the paid tier", then banned upselling inside a public or Pre-ETS channel. **Resolved:** the ban stands as the default, and **the funnel benefit is not counted anywhere in this plan.** The hybrid's paid tier is sold through B-family's own channel and carries B-family's selling cost. If a funder explicitly permits in-channel mention, that is upside, not a plan input. So "the hybrid is structurally best-shaped" (financial) is downgraded to a **cost-allocation observation:** the funder carries the shared fixed cost, and the family carries the marginal cost. It is not a conclusion that the hybrid is the best route.

Selling is labour, and it is counted (judge O2). A hand-sold family costs Sharon's time. Selling hours per paying family, s , equal the hours per prospect divided by the conversion rate. At 1 h per prospect and the probe's 3-in-10 pass rate, $s \approx 3.3 \text{ h}$ (~\$167 at $r = \$50$), **more than the price. Founder hand-selling therefore cannot be B-family's scaled channel.** It is only the test channel. B-family at scale needs a channel where Sharon's selling time is $\sim 0.5 \text{ h}$ per paying family, meaning referral-led or intermediary-led, and that is unproven. The probe records s and whether any paying family arrived through a low-touch referral (§6, G1).

Channel privacy constraints:

- No recruitment through the real student's school, district or personal network.
- Demos use the synthetic student only.
- No channel collects minors' disability data except the capped Stage 1 probe under the explicit exception in §6.

4.2 Launch sequencing

1. **Stage 0: preconditions** (this week to about mid-Oct 2026). Nothing external happens before these are done.
2. **Stage 1: three parallel tests** (about Oct 2026 to Jan 2027). The family probe fits the school calendar: current juniors (class of 2028) sit at the start of the Oct 2026 to May 2028 span the tool already models.
3. **G1** (target Feb 2027) picks **at most one** route.
4. **Stages 2-3: rebuild and route MVP** in one state. Target: ready by **Aug 2027**, for the **class of 2029 juniors**.
5. **Stage 4: first cohort plus the first aid-cycle re-verification** (fall 2027 to spring 2028).
6. **Stage 5: moat work, then possibly a second state.** VR/Pre-ETS allowability comes last and only if Sharon wants a state channel.

Resolution 5: school-year timing (judge tension 7). The rebuild (E-core 240-480 h plus a route add-on of 80-320 h, so 320-800 h) is 8-20 weeks full-time. At an unknown part-time H it is 6-12+ months. **Resolved:**

- The **probe families (class of 2028) are never the Stage 3 cohort.** They are served on their private concierge deploys through May 2028 under the commitment in §6, and they migrate only if the Stage 2 rebuild ships and they consent.
- **The Stage 3 product's first cohort is the class of 2029 juniors (fall 2027).** Working back from Aug 2027 gives a ~6-month build window after a Feb 2027 G1. That needs roughly **12-30 h/week** of build time.

- If H is lower, or G1 slips past ~March 2027, the first cohort becomes class of 2029 juniors starting mid-year (a shorter runway) or the class of 2030. State that plainly at G1, and do not rush the privacy rebuild to meet a school date.
- The fact-base re-verification is scheduled **before each Oct 1 FAFSA opening**.

4.3 Handed to team-gtm (D-003)

- Probe recruitment, screening and offer copy, within §4.1's privacy constraints and §6's probe exception
- The synthetic-student persona, demo script and assets
- The rewritten external pitch and per-segment messaging, bound by §3.2's banned claims
- Referral-partner mechanics (parent networks, consultants), including whether and how referrers are paid (uncosted today)
- The coordinator conversation guide, including how to elicit a price against the §6 pass mark
- Funder-facing narrative (budget figures come from §5.4)
- Conference, association and state-network presence, if B-district or (D) passes
- The one-state pilot playbook and the second-state playbook
- Content, SEO and app-store work: not recommended at plan level

5. Financial model

All \$ figures are dossier prices or arithmetic on stated parameters. None is a forecast, and none rests on a self-reported metric (D-005). **Soft assumptions are tagged [SOFT]**.

5.1 Parameters

Symbol	Meaning	Value used	Status
r	Loaded labour rate	\$50/hr	[SOFT] Upstream parameter, not a finding. Every hour-break-even scales with 1/r.
m	Per-student authored-layer hours, measured as LLM draft plus human review	0.25 / 0.5 / 1 / 3 h	[SOFT] The flip variable. Zero datapoints (R5). Stage 1 measures it.
s	Selling hours per <i>paying</i> customer	0.25 / 0.5 / 1 / 3.3 h	[SOFT] Zero datapoints. Newly modelled per judge O2.
F	Per-state fact-base build	40 / 120 / 300 h	[SOFT] R5
U	Per-state re-verification per aid cycle	≈ F/3 per year	[SOFT] Assumption
E	Build: E-core 240-480 h; + family 80-160 h; + district 160-320 h	12 – 24K; +4-8K; +\$8-16K	[SOFT] Low confidence (R10)
Legal, insurance	Privacy policy, terms, DPA; cyber/E&O	\$2-5K one-time; \$1-3K/yr	[SOFT] Unsourced placeholders; replace with quotes
Hosting	Platform	~\$600-1,800/yr	[SOFT] List prices; immaterial
Amortization	E, F and legal amortized over 3 years, straight line	—	[SOFT] Assumption

Symbol	Meaning	Value used	Status
Cohort	IEP juniors per state ~2.5-7.5K; IEP 11th-12th ~5-15K	—	[SOFT] Derived; Gap 16 unfilled
Renewal	One paid year per student; no renewal modelled	—	Conservative; renewal unknown

Fixed pools per state per year (F-mid):

- per-state content: ~4.0K(F – low 1.4K; F-high ~\$10K)
- platform, family variant: ~\$7-17K
- platform, district variant: ~\$9-20K
- platform, D variant: ~\$6-12K

5.2 B-family: family pays, per student, annual

Price: test at **\$99/yr per student**, with \$79 as the lower test point, \$49.99 as the floor never to go below, and \$149 as an upper probe.

- [SOFT] The price is set from cost, not from evidence (judge O3). The only anchor is parent-side IEP apps (~\$80/yr), and those are used every year of the IEP.
- **\$99 is 2x per child what CollegeCountdown charges for up to four children**, and CollegeCountdown also has a free tier. The premium has to come from the spine, promise aid and trade parity.
- Billing is annual. The tool spans ~20 months, and the second year's value is not modelled.

Contribution per paying family, before fixed costs: $c = P - (m + s) \cdot r - \sim 3.27 \text{ processing} - 1 \text{ LLM}$

At P = \$99	c	Break-even families/yr (fixed pool \$11-21K)	Share of IEP juniors (2.5-7.5K)	Sharon-capacity cap at H = 480 h, U = 40 h
m 0.25 + s 0.25	~\$70	~160-300	~2-12%	~880
m 0.5 + s 0.5	~\$45	~245-470	~3-19%	~440 (inside the break-even range)
m 0.25 + s 1	~\$32	~340-650	~5-26%	~350 (at the floor of break-even)
m 1 + s 1	~-\$5	never	—	—
s 3.3 (founder hand-sold, 30% conversion), any m	≤ ~-\$70	never	—	—
m 3 (hand-authored), s 0	~-\$55	never	—	—

Arithmetic correction: financial §4.1 printed "–\$105" for $99atm = 3h$. The correct figure is –55, as the judge noted. No conclusion changes.

Read. B-family breaks even only when $m + s \leq 1 \text{ h}$, and comfortably only at $m + s \leq 0.5 \text{ h}$. In that case it needs roughly **2-12% of a state's IEP juniors paying \$99 every year**. The **Sharon-capacity cap** sits at or inside the break-even range for anything above $m + s \approx 0.5 \text{ h}$, unless a paid reviewer at $r \approx \$25$ takes the

review, which doubles every m*. Raising the price to \$149 buys about 1 h of headroom, with no evidence families would pay it. The old "1-10% of the cohort" line is withdrawn. It divided by 11th-12th and assumed CAC = \$0.

5.3 B-district: per-IEP-11th-12th-student license

Price hypothesis: 15 – 20perIEP11th – 12thstudentperyear ** ,witha ** 3K/yr district minimum, on a multi-year term.

- [SOFT] **Unanchored above 5. * TAGG(5/set) and the top of the CCR band are the only anchors. Trinity's price (GP10) is the missing one.
- The minimum comes from $15 \times \text{an illustrative } 200 \text{ students} [\text{SOFT}, \text{GP8}] \text{ and replaces financial's } 2\text{K}$, which borrowed SameGoal's open (CT7) minimum and sits below the break-even price.

Economics (B-d2, the default: staff author, vendor support ~2-3 min per student):

- **Per student at \$15:** $c \approx \$12$.
- **Per 200-student district:** $\sim 2.5\text{K/yr}$ against 20 – 40h(1-2K) to sell, secure and onboard in year 1. **Year 1 is roughly break-even, and profit depends on renewal** (sticky per H7 for CCR, unverified for this buyer).
- **Per state:** with a fixed pool of \$13-24K, break-even needs **~1,050-1,950 students, i.e. ~5-10 districts of 200**. Each district costs 20-40 h to win.
- **B-d1** (vendor authors, $m \approx 5$ min) at \$15: $c \approx \$10.3$, so $\sim 1,260\text{-}2,330$ students. It is dominated by B-d2.
- **At \$10**, B-d1 needs $\sim 12\text{-}23$ districts and B-d2 $\sim 9\text{-}16$. Financial calls that "not plausible", and this plan agrees.

Read. B-district needs $\geq \$15$ plus renewal, or a single state-level buyer. Utah's contracts pre-empt a state buyer there.

- B-d2 moves the per-student labour onto case managers. That probably lowers WTP, and staff time may be the scarce input (the caseload caveat).
- **Neither assumption has been tested.** The coordinator test asks about it directly (\$6, M1.5).

5.4 D: funder pays; free to families and staff

Resolution 6: D is costed at its templated definition (judge tension 1). Financial sized the grant with the narrative m (24 – 167perstudent). That contradicts the D definition. **Corrected : ** D's per – student cost is intake support plus hosting, 1-5 per student [SOFT].

Grant ask, per state per year, per 1,000 students served, F-mid	Amount
Platform (D variant, E amortized)	~\$6-12K
Per-state content (F/3 + U)	~4K(F – high: 10K)
Per-student (intake support, ~\$1-5 each)	~\$1-5K
Total	**~11 – 21K/yr * (F – high: 17-27K)

Grant ask, per state per year, per 1,000 students served, F-mid	Amount
If an upkeep partner does F/U in kind	~\$7-17K/yr

- **Front-loading:** a funder may need to pay the build up front rather than amortized. A first-year ask is then ~\$25-45K [SOFT]: E-core plus a D intake plus the initial F, legal setup and a year of running costs.
- **This is the one number the funder scan (M1.6) takes into conversations.**
- The Pre-ETS-pot ratio (\$1.7-3.6K per served student) is **context only**. It is never a pitch figure and never evidence of a funding source (CT4, F8, H10).

5.5 Hybrid: D core plus B-family paid tier

- **Rule: the hybrid must be fundable at a paid-tier conversion of zero.** Its core is the \$5.4 ask.
- The paid tier (narrative and fit) is priced \$79-99. It is sold through B-family's own channel and carries B-family's m + s. Its contribution is upside only.
- **Do not project paid-tier conversion from the probe.** The probe has no free core beside it, so it overstates hybrid conversion (judge O3).
- Grant-term compatibility, upkeep scope and anti-supplanting are open (\$7, H1-H3).

5.6 Side by side

	B-family	B-district (B-d2)	D (templated)	Hybrid
Payer	Family	District special ed	Funder	Funder (core); family (optional tier)
Price / ask	\$99/yr per student (floor \$49.99)	15 – 20perIEP11 – 12student; 3K minimum	~11 – 21K/yrperstateper1,000served(7-17K with in-kind upkeep); ~\$25-45K in year 1 if the build is front-loaded	D ask; tier \$79-99
Closes when	m + s ≤ 0.5-1 h and ~2-12% of IEP juniors pay	≥ \$15, 5-10 districts per state, renewals	A funder and an upkeep owner exist	D closes on its own
Flip variables	m, s, WTP	Price, WTP, staff-time acceptance	Funder existence	Funder existence, grant terms
Evidence today	None (R1)	Contested (R2)	None (R3)	None (R1 + R3)

5.7 Stage 0-1 cost (what the tests themselves cost)

[SOFT] About **60-150 h of Sharon's time plus <\$1K cash:**

- synthetic student, including new-school fact research
- 10 family offers (~1 h each)
- 3-10 paid plans (m plus new facts)
- 3-5 coordinator conversations plus a Trinity demo
- a funder scan of 20-40 h

- up to ~\$1K in refunds or ~15-50 h of continuation support for probe families (§6)

6. Milestones

Each milestone names what unlocks it. **"Tests" replaces "retires"** for R1, R2, R7 and R8 (judge): a warm-referral probe tests R1 in concierge form and does not retire it for the product.

Stage 0: Preconditions (this week to ~mid-Oct 2026). Nothing external before these.

#	Milestone	Unlocked by	Unlocks
0.1	Live privacy stopgaps confirmed on the live site: (a) Netlify password protection or takedown; (b) <code>noindex</code> meta plus <code>robots.txt</code> <code>Disallow: /</code> ; (c) a check for search-engine or archive copies, with removal requests. This plan does not know whether they are in place (brief OQ2); intake and this lead did not re-check the site.	Sharon (minutes)	Everything else
0.2	The family is told about the multi-device data-loss risk (single-device rule), or sync is fixed	Sharon	Continued use of the one student's tool (a family matter, not a product decision)
0.3	Ask Sharon OQ4: hours split into fact base vs narrative; whether a structured-data or LLM pipeline produced the ~15 maps; her weekly available hours H	Sharon's answer	Sets expected m and the §4.2 build timeline; sizes Stage 1
0.4	Confirm the concierge offer may run purely as a B-family WTP test , not as a (C) business (brief OQ3)	Sharon's answer	M1.4
0.5	Confirm the go-ahead date (2026-09-26 vs 2026-09-25)	Orchestrator	Provenance only
0.6	Rewrite the external pitch against §3.2's banned claims	team-gtm / Sharon	Any external conversation
0.7	Hands-on Transition Elevated check (a sample draft plan), and open the SSIP PDF behind gtm's "worst place to start" call	Desk work	Any Utah conversation

Stage 1: Evidence tests (~Oct 2026 to Jan 2027; no product build)

#	Milestone	Route	Tests	Pass mark
1.1	OQ4 answer recorded (0.3)	All	R5 (Sharon's past hours)	A number or range per layer
1.2	Synthetic student built with a different district, school mix and disability category. Narrative drafted by LLM, then the review timed. Hours logged in four buckets: new-school facts (F expansion), category-specific content, LLM draft plus review (m), and fixes.	All	R5; the diagnosis generality gap (R11)	Measured m_AI+review; hours per new school; category content hours
1.3	Transition Elevated check (0.7)	Utah	R6	Utah stays the family-test state, and the district pitch there is confirmed or dropped

#	Milestone	Route	Tests	Pass mark
1.4	Family priced offer, n = 10 offers, Utah. See the design below.	B-family, hybrid tier	R1 (concierge form), R5, s, R7 (partial)	See G1
1.5	Coordinator budget-line test: 3-5 coordinators, mostly outside Utah. Synthetic demo plus caseload mock. Elicit a price. Also ask whether case managers would author (B-d2). Get Trinity's price and a demo.	B-district	R2 (tests CT6), R8 (partial)	See G1
1.6	Funder and upkeep-owner scan using the \$5.4 ask. In the same conversations, ask the hybrid questions: is a paid tier allowed? Does upkeep cover content the paid tier also uses? Anti-supplanting?	D, hybrid	R3, H1-H3	See G1

M1.4 design (resolving judge O1, O2, O3, O5 and O6 and tensions 4-5):

1. **n is fixed at 10 offers.** The verdict's bar is "3 or more of 10". The specialists' "3+ of 5-10" allowed a 3/5 pass, which is a different bar. Five offers cannot pass. **If 10 qualified families cannot be recruited through intermediaries by ~Dec 2026, the result is "inconclusive", and that is itself R7 evidence** (no expressed demand).
2. **Pay first, then collect data.** Families see **only the synthetic demo** and pay, or place a refundable deposit, **before** any of their own data is collected. The WTP signal needs no minor's data, and data is collected only from families who paid.
3. **The offer describes the Stage 3 product, not Sharon's labour:** "a planner with an AI-drafted plan reviewed by a person". This narrows (does not remove) the gap between concierge WTP and product WTP (judge O3). Paid-tier conversion is **not** carried into hybrid projections.
4. **The narrative is LLM-drafted with a timed review.** Hand-authoring is used only for a family that declines LLM processing, and it is logged separately as m_hand. That tests the m the business depends on (judge O1).
5. **Hours are logged by layer for every plan:** new-fact research (F expansion), LLM draft plus review (m), selling, including every non-converting conversation (s), and support. **s is computed per paying family.** Record how each paying family arrived (direct vs referral).
6. **Probe privacy exception.** See below.
7. **Terms for paying families if the product stops.** See below.

Probe privacy exception (resolving judge O5). The plan's rule is that no minors' disability data is collected before the Stage 2 rebuild. **The one stated exception is M1.4, capped at 10 families, under this minimum bar.** If any item cannot be met, fall back to a **synthetic-only** version: a WTP deposit on the synthetic demo, plan delivery deferred until after Stage 2, and m measured on synthetic students only (a weaker test).

- **Written guardian consent and student assent** name each field collected, where it is stored, who sees it, the LLM processing, the deletion date and the right to revoke.
- **Minimized intake:** only fields the schema needs. Diagnosis only if the family chooses to give it. No GPA required.
- **LLM processing** only on an API tier with no training and no retention, with names and identifiers replaced by placeholders before drafting. A family can decline, and its plan is then hand-authored.

- **Intake notes stored in one encrypted location** outside Supabase, the repo and email. Deleted within 30 days of the family's end date or on request.
- **A separate site per family**, each with its own password, `noindex` and `robots.txt`. No shared password. Authored content is never publicly readable. Never built by editing the real student's deploy.
- **Addresses and transit are not synced** (local-only or dropped). The single-device rule is in the consent.
- **No recruitment through the real student's school, district or network.**
- **No 11th family and no second round of 10** until either the exception is re-affirmed at G1 or Stage 2 ships.

Obligations to families who paid if the product stops (resolving judge O6). The terms are stated in the **offer**, before payment: "This is a test offer. If it is not continued, you choose one of two things. (a) Access continues through May of senior year, with one aid-cycle fact check and email support. (b) A full refund plus a PDF export of your plan." Budget: $\leq \sim \$1K$ of refunds, or $\sim 15\text{-}50$ h of continuation (one re-verification of these families' schools in summer 2027, plus $\sim 1\text{-}2$ h per family per year of support). **This budget is reserved before M1.4 starts.** Data is deleted per the consent on either path.

Gate G1: route selection (target Feb 2027). One reconciled definition.

Resolution 7: gate thresholds (judge tension 4 and O4). The offering gate let the hybrid proceed on a "borderline" family result. The gtm gate required a pass and omitted "both pass". Both used a coordinator bar of " $> \$0$ ". **Replaced by this single gate.** Every economic pass mark sits at the \$5 break-even.

Route	Passes G1 only if ALL of these hold
B-family	(a) ≥ 3 of 10 offers paid at $\geq \$79$. (b) Median m_AI+review ≤ 0.5 h, with no plan over 1 h. (c) At least one paying family arrived through a referral or intermediary with $\cdot 0.5$ h of Sharon's selling time. This is <i>signal</i> that a low-s channel exists [SOFT]; the probe cannot measure scaled s. 2 of 10 means one more round of 10, not proceed. 0-1 means stop.
B-district	(a) ≥ 1 coordinator names a real budget line and (b) states a price $\geq 15\text{per}EP11\text{th} - 12\text{thstudent}$, or $\geq 3K/\text{yr}$ for their district . (c) Trinity is judged a complement, not a copy. A named line at \$0-15 records CT6 as answered but fails B-district's economics , which then stops unless a state-level buyer surfaces.
D	(a) One grant or funding program and (b) one organization that commits to per-state upkeep, and (c) a program scale that can meet the \$5.4 ask ($\sim 11 - 21K/\text{yrperstateper}1,000\text{served}$, or 25-45K in year 1). Missing (a) or (b) makes D DON'T BUILD (verdict §7 step 5).
Hybrid	D passes, and B-family passes in full (no "borderline"), and the funder confirms in writing that a paid tier is allowed. Fundable at zero paid-tier revenue (§5.5).

Outcomes:

- **Exactly one route passes:** build it (Stages 2-3).
- **Several pass:** pick **one** for Stages 2-3. Never build two at once. The hybrid counts as one route. Tiebreak: the route whose pilot state has an existing fact base (Utah), then the lowest-m route.
- **B-district is the only pass:** it is a new product. Send its scope back to bizplan-offering. Its pilot state is where the passing coordinator is. That probably means a **new F at F-high**, with Utah's fact base sunk for this route. There is no parallel B-family Utah launch, and that removes the single-launch-state conflict (judge tension 6).

- **Nothing passes:** stop productizing. Honour the probe terms. The one student's tool continues as a family matter only.
- **Timing check at G1:** state which cohort Stage 3 can realistically reach, given H (§4.2).

Stages 2-5 (only after G1)

Stage	Milestone	Unlocked by	Unlocks	Unverified assumption it rests on
2	Privacy/tenancy rebuild, then content store with provenance, then sync plus timestamped events, then tests/repo. Raster art dropped.	G1 pass	Onboarding a real student beyond the probe	E estimate (R10, Low confidence)
3	Route-specific MVP (§2.3), one state. Target Aug 2027.	Stage 2	A sellable or fundable product	E add-on estimates; H (R5)
4	First cohort (class of 2029 juniors); one full aid-cycle re-verification before Oct 1, 2027 ; timestamped milestone data	Stage 3; a payer or funder contract	Real U, m and s, and first usage evidence (self-reported figures stay non-load-bearing, D-005)	WTP holds beyond the probe (R1/R2/R3)
5	Grow the trade/cert accommodations dataset and the fact base. Second state only after Stage 4 has priced a full upkeep cycle and Gap 16 is filled. VR/Pre-ETS allowability last, only if Sharon wants it.	Stage 4 results	Moat; a second state	F8 untested; R8

7. Risks & open evidence gaps

7.1 Inherited gaps (all ten carried, plus one the judge found under-carried)

#	Risk	What would close it	Where in plan
R1	Family WTP unknown; CollegeCountdown's free tier and \$49.99/4-children Plus	M1.4 tests the concierge form. Only Stage 4 retention and conversion beyond warm referrals retires it for the product.	M1.4, G1, Stage 4
R2	District WTP contested; \$0 default; Trinity's price unknown	M1.5 with the $\geq \$15$ pass mark, plus Trinity's price	M1.5, G1
R3	No funder, no upkeep owner	M1.6: both found, at the \$5.4 scale	M1.6, G1
R4	No VR/Pre-ETS planner precedent; allowability (F8) untested; Pre-ETS money not idle	RSA or state-VR allowability answer. Deliberately last.	Stage 5 (optional)
R5	Authoring hours unknown; no pipeline known	0.3 (Sharon), M1.2 and M1.4 m_AI+review logs by layer	Stage 0-1
R6	Utah statewide YouScience plus free Transition Elevated; \$9.33 not a price	0.7 hands-on check; coordinator test weighted outside Utah	0.7, M1.3, M1.5

#	Risk	What would close it	Where in plan
R7	No expressed demand; n=1 usage, no timestamps	<i>Partially tested</i> by whether 10 families can be recruited at all; retired only by Stage 4 timestamped use	M1.4, Stage 4
R8	Thin, closable moat	Trinity demo shows complement vs copy (partial). Retired only by a large, maintained fact base and accommodations dataset. Structurally thin, and may never retire.	M1.5, Stage 5
R9	Privacy rebuild for minors' disability data; FERPA for districts	Stage 2 item 1; the probe exception in the meantime	§6, Stage 2
R10	Engineering estimate Low confidence; content ops is the durable cost	Stage 2 actuals; Stage 4 measures one real upkeep cycle	Stage 2, Stage 4
R11	Diagnosis/category generality. Content (scholarships, accommodations framing, VR eligibility) varies by disability category, and F has no per-category multiplier. The schema's generality is unproven (verdict §6.4).	M1.2 uses a different category and logs category-specific hours. M1.4 families span categories.	M1.2, M1.4

7.2 Judge's objections: disposition

Obj.	Status after this plan	What would close what remains
O1 Wrong m measured	Resolved in design: LLM draft plus timed review, logs by layer, synthetic student separates F expansion from m. <i>The outcome is unknown:</i> m_AI+review may still exceed 0.5 h.	M1.2, M1.4 logs
O2 Hand-selling is uncostered	Costed; survives as the strongest structural objection. Founder hand-selling at 30% conversion costs more than the price. B-family needs an unproven low-touch channel.	G1(c) signal; Stage 4 channel data. Referral-partner economics (paying consultants?) are uncostered and handed to team-gtm.
O3 Probe WTP ≠ product WTP; no free core in the probe	Mitigated, not resolved. The offer describes the product, not Sharon's labour, and probe conversion is not used for the hybrid. A pass is necessary, not sufficient .	Stage 4 self-serve conversion
O4 District bar below break-even	Resolved: pass mark ≥ 150 or $\geq 3K$ per district. B-d2 staff-time acceptance remains untested.	M1.5 asks it directly
O5 Probe collects minors' data before the rebuild	Resolved: a stated, capped exception, pay-before-data, and a synthetic-only fallback	Stage 2
O6 No plan for paying families	Resolved: terms in the offer; refund or continuation budget reserved	—
O7 The paid layer is uneconomic or table stakes	Survives as a strategic risk. The plan answers it structurally (hybrid fundable at zero paid tier; the moat is the fact base) but has no evidence the fact base becomes defensible	Stage 5 dataset growth; incumbents' roadmaps

Obj.	Status after this plan	What would close what remains
Tensions 1-8	1 → Resolution 6; 2 → Res. 3; 3 → Res. 4; 4, 5 → Res. 7 and n = 10; 6 → G1 "B-district only" rule; 7 → Res. 5; 8 → arithmetic corrected	—

7.3 Plan-specific risks

#	Risk	What would close it
H1	Grant terms forbid a paid tier (hybrid)	Written funder answer in M1.6
H2	The funder's upkeep obligation doesn't cover content the paid tier uses	Same conversation; the default design keeps all facts in the free core, so this only matters for the narrative layer
H3	Anti-supplanting / Pre-ETS positioning clashes with charging families (untested)	Funder or RSA answer. Default: no paid tier inside a Pre-ETS channel.
P1	LLM processing of real families' data (a third-party processor)	Consent plus a no-retention API tier plus placeholder substitution; a hand-authored fallback
P2	The synthetic student is inferable if it resembles the real student	Different district, school mix and category; reviewed before any demo
T1	H (Sharon's weekly hours) too low to hit Aug 2027	0.3 answer; timing check at G1; slip the cohort, never the privacy rebuild
C1	Referral partners (consultants at ~\$140/hr) have no reason to refer to a \$99 product unless paid	team-gtm designs; cost it before Stage 3
C2	Cohort shares rest on Gap 16 and district sizing on GP8	NCES headcount before any share is quoted; coordinator caseload data from M1.5
V1	Go-ahead relayed, not seen; date off by one	0.5
L1	Live exposure of the student's details (verdict \$0) is unconfirmed as fixed	0.1

7.4 What this plan cannot resolve (stated plainly)

- **Whether anyone pays.** Every route's revenue line rests on evidence that does not exist yet.
- **Whether a low-selling-cost family channel exists.** The probe measures founder selling, not scaled selling. G1(c) is only a signal.
- **Whether the paid narrative has durable value** once an incumbent adds AI narratives (O7).
- **Whether a funder will accept a paid tier (H1-H3).** The plan only makes the hybrid survivable without one.
- **Whether case managers will author plans (B-d2)** under caseload pressure.
- **Sharon's hours and pipeline (OQ4).** Asked for in Stage 0, not yet answered.

8. Plan confidence

Overall: Low as a business plan. Moderate as a Stage 0-1 evidence plan.

- **Test plan: moderate.** The judge rated the test sequence "earned": verdict-first, privacy first, VR last, and honest about pricing. This plan applies all seven of the judge's required fixes:
 1. m measured as AI draft plus review, logged by layer
 2. a district price threshold at break-even
 3. an explicit probe privacy exception with a synthetic-only fallback
 4. D costed at its templated definition
 5. one reconciled G1 with $n = 10$
 6. selling costed as hours
 7. terms and a budget for paying families

With those in place, G1 can no longer return "go" on evidence that does not support the build. That was the judge's central complaint. The test plan does not rise above moderate because 0.1 (live privacy) and 0.3 (hours) are unconfirmed, and the probe still overstates product WTP (O3).

- **Economics: low.** No route closes on present evidence:
 - B-family needs $m + s \leq 0.5 - 1 \text{ h}$ **and** a few percent of a state's IEP juniors paying \$99 a year, with founder selling already shown to cost more than the price.
 - B-district needs an unanchored $\geq \$15$ price across 5-10 districts per state.
 - D needs a funder that the dossier has no evidence exists.
 - The hybrid needs D **plus** grant terms nobody has asked about.

The judge put business-plan confidence at low. The fixes make the plan honest; they do not make the economics better, and several make them visibly worse (O2).

- **What would raise it.** A funder plus upkeep owner at the \$5.4 scale would raise D and the hybrid to moderate. $m_{\text{AI}+\text{review}} \leq 0.5 \text{ h}$ plus a 3/10 pass plus a referral-sourced payer would raise B-family to low-moderate. A coordinator quoting $\geq \$15$ would do the same for B-district.
- **Consistency with the verdict.** The plan builds nothing before the verdict's §7 tests report. It keeps (A) excluded. It uses (C) only as a test. It carries R1-R10 plus R11. It stays at or below the verdict's confidence for each intent.

9. Prior-run cross-reference

This is the first plan for this target. The team-business-plan run log has no entry for after-graduation-planner. The two prior entries (the claude-code-agent-monitor SaaS reframe and ai-teaching-patterns-book) are unrelated targets. A pattern they share with this plan: both were built on conditional or NEEDS MORE EVIDENCE verdicts, and both put an evidence gate ahead of any cash or build. This plan does the same with G1.

The verdict itself is also a first evaluation (verdict §8). There is no revised verdict or changed MVP to reconcile.

Verdict: after-graduation-planner

Prepared by viability-lead, 2026-09-25. This synthesizes `brief.md`, `supporting/product.md`, `supporting/market.md`, `supporting/stakeholder.md` and `judgment.md`. It was cross-referenced against `~/claude/skills/team-viability/memory/run-log.md`, which has no prior entry for this target (see §8).

INTERNAL FILE. The prototype belongs to a real minor, called "the student" here. Nothing in this file may be quoted in external output or used in an outside query (D-003).

ACT NOW, INDEPENDENT OF VERDICT: the student's private details are public

This is not part of the BUILD / DON'T BUILD question. It is live harm to a real minor, and it should be fixed before anything else in this file.

- **What is exposed.** The product specialist found that all authored content ships as static constants in the public HTML at [the live site](#). The orchestrator confirmed this independently with a keyless `curl`. **Anyone with the URL can read the following without a key:**
 - the student's first name
 - GPA
 - high school and district
 - diagnosis (dozens of mentions, including a diagnosis-specific scholarship card)
 - a first-person narrative about the student's interests and activities
 - household structure, including the three homes, per-household transit times and a relative's business
- **Why crawlers can reach it.** The page has no `noindex`, `/robots.txt` returns 404, and Netlify serves the page with `cache-control: public`.
- **What the key actually protects.** The link key protects saved progress only. The brief's claim that per-student data needs the key is wrong for everything except progress.
- **Family addresses.** These are user-entered, not in the code. [Storage details removed for privacy.]

Stopgaps. Each takes minutes. They are listed from most to least protective:

1. Turn on Netlify password protection for the site, or take it down until the authored content sits behind the key.
2. At minimum, add `<meta name="robots" content="noindex">` and a `robots.txt` with `Disallow: /`.
3. Afterwards, check whether search engines or archives have indexed the page, and request removal if they have. The stopgaps limit future exposure only. They do not undo anything already crawled or cached.

Second live risk (data loss, not privacy). `save_progress` sends the whole snapshot, and the last write wins. On load, local keys that are missing from the remote copy get deleted. If the student and a parent use different devices, one can silently erase the other's work. The tool is meant to be used until May 2028, so tell the family to use one device, or fix the sync logic.

No agent in this run changed the site or the backend. This is a report only.

1. Target summary

"After Graduation" is a gamified transition planner Sharon hand-built for one real high-school junior with an IEP. It is framed as a TV show (Season 1 is high school, Season 2 is after graduation) and has three Episodes:

1. **Explore Schools.** A drag-to-rank race track of candidate schools. Each school has a modal covering cost/aid, credential-to-salary outcomes, disability accommodations, personalized reasons and reflection prompts.
2. **Financial Planning.** A GPA tile and flip-card decks for VR, Pell, state promise aid and disability scholarships.
3. **Important Dates.** A dated milestone timeline from Oct 2026 to May 2028 that mixes IEP, transition-plan and Summary of Performance steps with school, career and money deadlines.

Plain-language copy runs throughout, and trade and certificate paths get the same weight as 4-year college. The site is one static HTML file on Netlify. Supabase RPC persistence is keyed by a link fragment, and there is no login. All content is hand-authored for one student, district, state and diagnosis. The question is whether this one-off should become a product.

2. Verdict

NEEDS MORE EVIDENCE. This is the top line. It is split by Sharon's intent, as the judge recommended, because the evidence splits by intent.

Component (Sharon's intent)	Status	Confidence
(A) Venture-scale product or self-serve SaaS	DON'T BUILD	Medium-high
(B) Modest paid tool: family subscription or district special-ed license	NEEDS MORE EVIDENCE	Moderate
(C) Done-for-you service or consulting, with the planner as the deliverable (Utah families first)	NEEDS MORE EVIDENCE (cheapest to resolve)	Moderate
(D) Free, nonprofit or grant-funded offering	NEEDS MORE EVIDENCE	Low-moderate

What each component means:

- **(A) DON'T BUILD, on the evidence already in hand.**
 - IEP-only revenue at district CCR prices is roughly \$1.5-4M a year at 100% national share. This is the market specialist's illustrative figure, but its direction survives a 2x correction.

- Every payer is weak or empty.
- The implied launch state is already pre-empted by a statewide CCR contract and a free state-listed transition tool.
- Edtech capital has contracted, and a raise on n=1 with no demand signal would not land.
- The judge's point stands: no cheap evidence could turn this into a venture case, so calling it **NEEDS MORE EVIDENCE** would be a cop-out.
- **(B) and (C) are genuinely open.** Specific, cheap evidence could flip them either way (§7).
 - (C) is the path the stakeholder specialist rated most plausible, and it is also the cheapest test. It can be sold by hand to families with **no product build**, beyond a private per-family deploy with the privacy fixes above.
- **(D) is open.** Families would not pay, so the questions are whether a funder exists and whether per-state content upkeep can be staffed. The dossier has no evidence on nonprofit or grant buyers (GP7).

Not a productization verdict. Continuing to run the existing tool for the one student is a family matter, not a product decision. It has no row in the table above so the bizplan-triage gate cannot treat it as a BUILD component. It should continue only after the §0 stopgaps.

Gate for team-business-plan. Do **not** start team-business-plan yet. Under the bizplan-triage gate this is HALT-PENDING-GO-AHEAD for (B)-(D) and HALT-PENDING-OVERRIDE for (A). It should stay halted until Sharon answers the two open inputs below and the concierge test in §7 reports back. Planning while Sharon's intent is unknown and there is no payer evidence would be wasted work.

Open inputs that would finalize this verdict. Both are questions for Sharon, not research gaps:

1. **Sharon's intent (brief OQ7).** Which of (A)-(D) is she after? If the answer is (A), the verdict becomes a final DON'T BUILD with no further testing. Any other answer narrows the verdict to one component and one test.
2. **Sharon's authoring hours (brief OQ4), split two ways, and whether a structured-data or LLM pipeline already exists.**
 - The split: hours spent on **Utah and school fact research** (reusable per state) versus hours spent on **this student's narrative and fit judgments** (per student).
 - The product specialist found about 15 content maps serialized in Python `json.dumps` style. That suggests structured source data exists somewhere outside the artifact.
 - This number decides whether the per-student marginal cost is minutes or days (see §6, tension 1).

3. Confidence

Moderate overall. Medium-high on the (A) DON'T BUILD. Moderate on (B) and (C). Low-moderate on (D).

- **Why not higher.** The two inputs that move the answer most are Sharon's intent and her hours. Neither has been asked yet. The judge noted that an unasked question is procedurally weaker than a real research gap, and I agree. The confidence stays moderate until those answers come in, not because the evidence is thin.
- **Where the judge moved it.** The judge's challenge had a real effect in two directions:

1. **It hardened (A) from "open" to DON'T BUILD.** Read together, the specialists leaned toward a single NEEDS MORE EVIDENCE. The judge showed that for a venture or SaaS intent, the market math and payer evidence already decide the answer. I ratify that split.
 2. **It weakened the stakeholder's cost objection without removing the payer objection.** The stakeholder priced the whole plan as per-student labour at \$50/hr (4-7 minutes per student at district prices). The product specialist's breakdown shows that most content is reusable per state or nationally. Only the narrative and fit layer is per-student, and that is the lowest-risk layer for an LLM to draft. So I downgrade the "authoring cost kills district and VR models" claim from High to Medium/Low until OQ4 is answered. **This does not raise the verdict to BUILD.** The bottleneck moves from "each plan costs too much" to "nobody has been shown to pay for plans", and all three specialists and the judge agree on that second point. Cheap AI personalization is also table stakes, because Trinity already does it for the district special-ed buyer.
- **What stays high.**
 - The need is real (H1, H11, primary sources).
 - No payer has shown willingness to pay (GP1, CT5, CT6, H10). This is an absence of evidence, and absence is all the evidence there is.
 - The privacy model needs a rebuild, not a patch.
 - The client-code findings. The product specialist read the actual code and confirmed it against the live deploy.
 - **What the judge lowered.**
 - "The prototype is in real use" is overstated. The evidence is one saved answer and two picks, with no timestamps. The student interacted with it at least once. Sustained use is unverified.
 - The multi-student MVP estimate of "weeks" moves to the high end of its range. Moving authored content server-side is an architecture change, because that content currently lives in static page constants.

4. Case for

- **The need is documented, and the design fits the real population.**
 - IEP families have documented information gaps on training options, aid, VR and self-advocacy (H11).
 - IEP youth are about twice as likely as other students to expect a sub-4-year path (H1). The prototype's trade and certificate parity and its IEP-dated timeline match how this population actually plans.
- **The feature fit with the remaining white space is strong.** The narrow unoccupied space is one timeline, dated by IEP milestones, that combines VR/Pre-ETS, aid (including state promise aid), trade and certificate paths and a per-student narrative. The prototype already does that.
 - CollegeCountdown covers roughly half of this for families.
 - Trinity covers roughly half for districts.
 - Neither advertises the merged dated spine.
- **The front end and UX are largely done, and they are good.**
 - UX finish is high.
 - Accessibility work is real: 93 `aria-*` attributes and reduced-motion handling throughout.
 - The content maps (about 15) already act as a working content schema.

- A one-state multi-student MVP is mostly back-end and content work. The UI would need little change.
- **Per-state amortization plus AI drafting could make per-student cost small.** This was the judge's strongest pro-BUILD point.
 - Utah aid, VR, school and glossary content spreads across every Utah student.
 - The truly per-student layer is judgment about fit, not dollar or deadline facts, so it is safer to have an LLM draft it.
- **Low Pre-ETS reach (6-10%) is a real gap** that a well-timed "apply to VR" milestone speaks to. The caveat is that the bottleneck may be caseworker capacity, which software does not fix.

5. Case against

The judge's strongest counter-argument is **DON'T BUILD for any venture or SaaS intent, and no payer has evidence of paying for any intent.**

- **Market size.** IEP-only revenue at district CCR prices (\$3-6 per student) is about \$1.5-4M a year even at 100% national share.
- **Payers.**
 - Families have a free incumbent (CollegeCountdown's free tier) and no willingness-to-pay data (CT5).
 - District CCR buyers are not buying IEP features.
 - District special ed defaults to \$0 for student-facing tools (CT6), and Trinity already sells there.
 - VR has never bought a planner (H10), and whether it is even allowed is untested (F8).
 - Ads- or data-funded models are ruled out for minors' disability data, given the Naviance settlement and the PowerSchool breach (H8).
- **Utah launch headwind.** Utah is the cheapest launch state because the content exists, but it may also be the most pre-empted. It already has a statewide CCR contract (YouScience) and a free USBE-listed transition tool (Transition Elevated). Per F2, the free tool is a headwind, not proof that price is capped.
- **Thin moat.** The white space is Medium confidence and judged from marketing pages only (GP10). The most distinctive element, the per-student narrative, is the hand-authored part. Either incumbent could add a VR milestone quickly.

This counter-argument is fully reflected in the verdict. It is why (A) is DON'T BUILD and why no component reaches BUILD.

6. Cross-specialist tensions

These are carried forward from the judge's §3. Each is reflected in the verdict or confidence above, not just noted.

1. **Per-student vs per-state cost (stakeholder vs product).** The stakeholder wrote without the product specialist's output and priced everything as per-student labour. Product splits the content into reusable-nationally, reusable-per-state and per-student layers. The two give very different break-evens.
 - Resolution: I adopt product's split, with the judge. The stakeholder's break-even table is a ceiling, not an estimate.
 - Sharon's OQ4 answer, split into fact-base hours and narrative hours, settles it.

2. **"Near-total rebuild" vs "weeks for one developer."** Both specialists agree the durable cost is content operations. The privacy finding tilts the engineering side toward the stakeholder: moving authored content behind auth is an architecture change. Product rated its own estimate Low confidence.
3. **The stakeholder's pitch overclaims against its own sources.**
 - Trinity *does* compare schools and filter by cost and disability support. The pitch is accurate only on dated deadlines and aid/VR funding.
 - "A real student is using it today" overstates n=1 interaction with no timestamps.
 - "Plugs the Pre-ETS leak" ignores the caseload caveat.
 - Any external pitch must be rewritten before use.
4. **"Feature fit" is the fit of a one-off.** The market specialist scored capability coverage from the brief's description. The product specialist shows how much of that coverage is tied to one student, diagnosis, relative's business and state. The fit for a generalizable product is unproven.
5. **The unique data asset is thin.** The only truly unique data is accommodations for trade and certificate schools (X2). That is currently a handful of records among five schools.
6. **Utah: cheapest to launch and possibly most pre-empted.** The market specialist flags this. The product and stakeholder specialists do not. It must be addressed before any Utah coordinator or VR conversation (§7).
7. **Privacy appears in only one lens.** The market and stakeholder specialists treat privacy as a future compliance cost. The confirmed live exposure shows the current threat model assumed the key protected the content. For a product, that confirms brief OQ8 as a rebuild cost. Separately, it is the "act now" item at the top of this file.

7. What evidence would resolve it

Listed from cheapest to most expensive. Steps 1 and 2 are questions for Sharon and should come back before anything else runs.

1. **Ask Sharon her intent (OQ7).** Pick one of (A)-(D).
 - (A) makes the verdict a final DON'T BUILD, and the run stops there.
 - Any other answer narrows the verdict to one component and the test below that matches it.
2. **Ask Sharon her hours (OQ4).**
 - Split them into Utah and school fact research versus this student's narrative and fit judgments.
 - Ask whether structured source data or an LLM/script pipeline produced the content maps.
 - Compare the per-student figure with the break-evens at \$50/hr:
 - family price (about \$80/yr): about 1.6 hours per student
 - district price (\$3-6 per student): minutes
 - consulting price (\$140+/hr): unbounded, because the hours are the product
3. **Concierge payment test, for (C), and for the family side of (B).** This takes weeks and needs no product build.
 - Offer a done-for-you plan at a real stated price to 5-10 Utah IEP families.
 - **Proposed resolving threshold:** 3 or more of 10 pay or put down a deposit. 3 or more means (C) moves to BUILD as a service. 0-1 means DON'T BUILD. 2 means run another round.
 - Log the authoring hours for each paid plan. That is the real OQ4 datapoint for a second student.

- **Privacy constraint:** demos must use a **synthetic student**, never the real student's plan or details (D-003).
4. **Coordinator budget-line test, for the district side of (B).**
 - Ask 3-5 transition coordinators which budget line they would pay from, and at what price per student or per site.
 - Resolving signal: at least one names a real budget line and a price above \$0. This tests CT6 directly.
 - Before any Utah coordinator talk, check what Transition Elevated actually covers. It is free and state-listed, and could make the Utah district pitch unnecessary.
 - Get Trinity's price and a demo (GP10) to tell whether a district product would be a complement to Trinity or a copy of it.
 5. **Funder signal, for (D).** Find one grant or foundation program that funds transition-planning tools, and one organization willing to staff yearly per-state content upkeep. Without both, (D) moves to DON'T BUILD.
 6. **Optional, supporting only.** With Sharon's explicit permission, a read-only look at the student's `load_progress` data (never `save_progress`). Even then it cannot show milestone *timing*, because checked steps are stored as "1" with no timestamps. It will not resolve the verdict on its own.

State VR and Pre-ETS allowability (F8) is **not** on this list. It is the slowest and least likely route, and it matters only if Sharon explicitly wants a state channel.

8. Prior-run cross-reference

This is the first evaluation. `~/claude/skills/team-viability/memory/run-log.md` has no entry for after-graduation-planner. The brief also records no prior `verdict.md`.

One process note: the judge dates the orchestrator's keyless re-check 2026-09-26, but this run is dated 2026-09-25. That is probably a typo in `judgment.md`. It does not change the finding.

Dossier: after-graduation-planner

Lead (team-research) | Run: 2026-09-25 | Pilot pass, 4 angles Inputs: `brief.md`, `comparison.md` (15 agreements / 16 gaps / 9 contradictions), `validation.md` (24 claims: 21 CONFIRMED / 0 UNSOURCED / 3 CONTESTED). References like "AG9" or "CT5" point to `comparison.md`. "Val" means `validation.md`.

Readers: team-viability (market specialist) and team-business-plan (financial specialist). Sections 2a-2c are lookup tables for competitors, price points and payers. Every row carries its source tier and verdict, so you can lift a row without re-checking it.

Privacy: this dossier holds only public product and program names and generic capability terms. It contains no student-identifying or unreleased-product specifics (D-003).

1. Topic summary

The idea is to turn a hand-built, one-student prototype into a product. The prototype is a personalized, game-like planner that helps a US high-school junior or senior plan what comes after graduation. It has four parts: a school explorer, a financial-aid deck, a dated timeline and a plain-language glossary. The school explorer ranks schools on cost, aid, credential-to-salary outcomes and disability accommodations, and gives reasons tailored to the student. The aid deck covers FAFSA/Pell, state promise aid, Vocational Rehabilitation and disability scholarships. The timeline interleaves IEP transition milestones (transition plan, Summary of Performance, VR referral) with application, aid and career deadlines. The product is aimed especially at students with IEPs and gives trade, certificate and technical-college paths the same weight as 4-year college. The research question: does anyone already do this, and could it be a viable product? The answer feeds team-viability's BUILD / DON'T BUILD / NEEDS MORE EVIDENCE call and, if that isn't DON'T BUILD, team-business-plan.

2. Headline findings (CONFIRMED only)

One or two per angle. Confidence is the combined rating after validation.

market-size

- **H1. The non-4-year path is the majority, and it is larger still among IEP youth.**
 - BLS (Apr 2025): 62.8% of 2024 graduates were enrolled in college that October, and about a third of those were at 2-year colleges. So about 58% of a class is not at a 4-year college that fall.
 - NLTS 2012 (Table ES8, read directly): among youth, 25% with IEPs vs 14% without expect postsecondary education short of a 4-year degree. Among parents, 27% vs 14%. That is roughly twice the rate.
 - Primary. **High.** (AG6)
- **H2. Pre-ETS reaches a small minority of eligible students, and schools make most referrals.**
 - Reach is about 180-300K students served, roughly 6-10% of those eligible. Eligibility covers students with an IEP or a 504 plan.

- Schools make 71-87% of referrals (NCLD 2026, p.15, read directly). Hechinger (Feb 2025) gives 295K of 3.1M.
- Two independent sources. **High on direction.** Give the size as a range, never a single figure (AG5, CT3).
- **H3. No public source sizes transition-specific spend.**
 - IDEA Part B is \$14.21B (FY24 and FY25, CRS R44624) and is not earmarked for transition.
 - No transition-software market figure exists, and the syndicated "special education software" reports aren't usable.
 - **High** that the figure can't be found in public sources (AG14).

competitors

- **H4. No CCR platform publicly offers an IEP transition module. The narrow white space is still open, but it is smaller than the comparator said.**
 - Xello, Naviance, SchoolLinks, Kuder, MajorClarity and YouScience document no IEP module. The validator checked independently.
 - Two products the scouts missed or under-described each cover roughly half the concept (Val falsification test):
 - **CollegeCountdown** (B2C family app, 2026): IEP/504 hub, CeDaR-powered disability-aware college search, 700+ disability scholarships, FAFSA checklist, deadline tracker, 44-milestone roadmap, XP/streak gamification.
 - **University Startups "Trinity"** (district special-ed, 2026): AI-generated IEP transition plans that recommend colleges and 5,000+ training programs, filtered by disability support and cost.
 - **What remains unoccupied:** one timeline that is *dated by IEP milestones* and combines VR/Pre-ETS, aid (including state promise aid), trade and certificate paths, and a per-student narrative. Neither product advertises VR/Pre-ETS, the Summary of Performance or IEP-meeting milestones.
 - **Medium.** Everything here comes from public pages and vendor marketing. No demos were done (GP10).
- **H5. Trade and certificate coverage on its own is not a differentiator.** YouScience, MajorClarity/Edmentum, Trinity, CareerVillage Coach and BigFuture already treat those paths as first-class. **High** (AG6).
- **H6. Free, publicly funded IEP transition tools already exist.**
 - iTransition, WiTransition and Transition Elevated each turn a short survey into a draft plan.
 - None of them compares schools, covers aid or tracks dated deadlines.
 - **High** (AG2). The claim that these free tools "cap price" was not confirmed. See section 5.

monetization

- **H7. District CCR sells at about \$3-6 per student per year on multi-year, sticky contracts, and more states now buy one platform for every district.**
 - Statewide purchasing is confirmed in WI, FL and UT, and in WA (all districts by 2026-27).
 - District quotes and contracts back the \$3-6 band (table 2b).
 - **High** on the band and the contract structure. **Low-medium** on total market size, which nobody publishes (AG9, GP3).

- **H8. Student-facing planning software is mostly free to students and paid for by institutions or recruiting marketplaces.**
 - Examples: Kollegio, CollegeVine Sage, Scoir's student side, Naviance/Intersect, CareerVillage Coach, BigFuture.
 - Privacy exposure makes an ads- or data-funded model risky for disability data: Naviance's \$17.25M settlement and the 50-60M-student PowerSchool breach.
 - **High** (AG11, AG12).

iep-transition-niche

- **H9. Pre-ETS is a ring-fenced pot of about \$0.5-0.65B a year, and it mostly buys human services.**
 - At least 15% of the federal VR grant is reserved. It can fund only the five Sec. 113 activities, and none of it can go to admin.
 - It is delivered mainly through provider and district contracts.
 - No RSA guidance addresses planning-software licenses.
 - A tool would have to be framed as delivering a *required* activity (job exploration or postsecondary counseling). "Authorized" activities are funded last.
 - **High** on the rules. Whether a planner is allowable is **untested inference** (AG3, AG4, GP6).
- **H10. VR money has bought student-facing software, but only career-exploration video, never a planner.**
 - Nevada DETR VR bought 19,978 Virtual Job Shadow licenses (every student with a disability aged 14+, 2018 to at least Oct 2022).
 - Massachusetts MRC funded about 40 Pre-ETS providers to buy that kind of software, and Minnesota funded provider subscriptions.
 - Every precedent is the same vendor, Virtual Job Shadow (now Pathful). No dollar amounts were found.
 - **Medium-high.** The state records were seen as search extracts (CT1).
- **H11. The need is documented: IEP families have real information gaps.**
 - The gaps cover training options, aid, VR and the self-advocacy that college accommodations require (GAO-24-105614; NLTS 2012).
 - **High** that the need exists. **Medium** on how big it is today, because the national parent data is from 2012-13 (AG7).
 - *Expressed demand for a tool* was **not** found (GP1).

2a. Competitor map (for the market specialist)

The capability letters come from the brief:

- (a) personalization beyond interest inventories
- (b) trade and certificate paths as first-class
- (c) IEP milestones merged with application, aid and career deadlines
- (d) per-school accommodations, VR and disability aid
- (e) game or story framing and plain language

"No evidence" means nothing public was found. It is not a verified "no" (GP10).

Product	Channel / payer	Overlap with concept	Missing (per public evidence)	Source / verdict
CollegeCountdown	B2C family, freemium ("free forever" plan; paid tiers not shown)	IEP/504 Transition Hub; CeDaR-powered "DS support tier" school filter; 700+ disability scholarships; SAT/ACT accommodations deadline calculator; FAFSA checklist; deadline tracker; 44-milestone 9th-grade-to-enrollment roadmap; XP, streaks and leaderboards (for test prep)	VR/Pre-ETS, Summary of Performance, IEP-meeting milestones, trade/certificate paths, state promise aid, story built on the student's own plan	Vendor marketing and App Store listing (2026-06-29) only. Found by the validator, not by the scouts. Review counts and outcome claims unverified.
University Startups "Trinity"	District special ed / state agency, paid; price not public (60-day money-back guarantee)	AI IEP transition-plan generation; College Agent (30K+ 4-year and community colleges); Training Agent (5,000+ vocational programs); filters by disability support, cost and location; Indicator 13 compliance dashboards	Financial aid/FAFSA, VR funding, any dated timeline	AWS ML blog (2026, co-written with the vendor), vendor site, logo wall (NYC transition centers, Poudre SD, Grand Prairie ISD, ND DPI). "More than a dozen states" in year one is self-reported.
Xello	District/state CCR, ~\$3.75-5.45 per student	Interest inventories; "The Real Game" (game-like in parts); statewide in WI and FL	No IEP module (the special-ed page is teaching guidance)	Help center; confirmed
Naviance (PowerSchool)	District CCR	College search, applications	No documented link to PowerSchool's own IEP system (Special Programs)	PowerSchool docs; confirmed
SchoolLinks	District/state CCR (statewide in WA)	CCR, applications	No IEP module found	Confirmed (absence)
Scoir	District CCR; free student channel	College search	No IEP module found	Confirmed (absence)
YouScience	District/state; ~\$49 family test	Aptitude "brain games"; 200+ certifications (b)	No IEP module found	CO 1.7; confirmed (absence)
MajorClarity (Edmentum, since Oct 2025)	District CTE	Trade/CTE paths (b)	No IEP module found	Edmentum release; confirmed
Kuder Navigator	District CCR	CCR	No IEP module found	Val; absence
iTransition / WiTransition / Transition Elevated	Free, publicly funded	Survey-to-draft transition plan	School comparison, aid, dated deadlines	DRC Kansas, App Store, USBE; confirmed

Product	Channel / payer	Overlap with concept	Missing (per public evidence)	Source / verdict
T-Folio, Next Steps NH, TWC Pre-ETS Catalog	Free, public	Transition portfolios and catalogs	Same as above	Scout-cited; not individually re-checked
TAGG (Zarrow)	District special ed, \$5.00 per set	Transition assessment	Not a planner	Official Zarrow page (corrects the older ~\$3)
ULS transition curriculum	Sold per provider	Transition curriculum	Not a planner	CO 2.4
Pathful Explore (formerly Virtual Job Shadow)	Bought by VR agencies (NV, MA, MN; GA weak)	Career-exploration video	Not a planner	CT1; confirmed
Kollegio	Free; seed-funded (\$3.55M total); no revenue model yet	AI college counselor	IEP, VR, trade focus	SiliconANGLE; confirmed
CollegeVine Sage	Free (college-recruiting marketplace); paid human mentoring ~\$1,300 on average	AI counselor	IEP, VR	Confirmed
CareerVillage Coach, BigFuture Credential Finder	Free, nonprofit	Career coaching; credentials (b)	IEP, VR	Confirmed
USC Collegeology / FutureBound	Research games (FutureBound targets middle school)	Game framing with self-efficacy evidence	Generic simulation, not the student's own plan; may be dormant (unverified)	Research-group reports and an IES-funded RCT
CeDaR; K&W Guide (\$35.99, 350+ colleges, LD/ADHD/ASD); Think College Search (300+ ID programs)	Free / book / free	Per-school disability-support data sources (d)	CeDaR covers degree-granting schools only, detail for "nearly 700"; trade schools only "planned"	Val AG8; possible data inputs as well as competitors

Consolidation signals:

- PowerSchool bought Hobsons/Naviance for \$318.9M (2021), and Bain took PowerSchool private for \$5.6B (closed Oct 2024).
- MajorClarity passed from Paper (the acquisition date is disputed: 2022 or Feb 2023) to Edmentum (Oct 2025).
- Edtech VC fell from \$16.7B (2021) to \$2.6B (2025, Tracxn).

- "Standalone CCR is hard to sustain" is inference, not a sourced finding (AG10).

2b. Price points (for the financial specialist)

Item	Price	Basis / caveat	Verdict
Xello, district HS	**4.45 * <i>paid</i> (5.45 regular), 3-year term	Park Hill SD board-posted quote, May 2025	Confirmed (Val)
Xello, district MS	3.75 <i>paid</i> (4.75 regular)	Same quote	Confirmed
Xello order form (MO A3)	\$4.95 list / \$4.45 paid	The validator could not re-find \$4.95; list prices vary by district and year	Paid figure matches; list figure not re-found
Scoir	\$4.80 average	Vendor comparison sheet	Vendor claim
Scoir, real contract	2.52(2024 – 25), 3.96 (2025-26)	One Toledo, OH school	Confirmed
SchoolLinks	\$5.51 + fees	A <i>competitor's</i> (Scoir's) comparison sheet	Weak
Naviance	\$3-15 per student; Seattle \$619,538 over 3 years	2017 sources	Confirmed but dated
Florida statewide Xello	~\$2.8M/yr ($\approx$ \$2.50 per active student, 1.1M+ users)	One advocacy-leaning source (The Drill Down); no FDOE or vendor figure	See section 5
Florida, Sept 2026	\$6.6M federal DOL/ED grant to expand Xello	New (Val)	Confirmed
Utah YouScience	~\$9.33 per student	\$3.5M budget <i>request</i> $\div$ ~375K students; derived, not a contract	See section 5
TAGG transition assessment	\$5.00 per set	Official Zarrow page	Confirmed (corrected)
SameGoal (IEP software)	1.80 – 2.65 <i>per student</i> ; <i>district minimums</i> 2.1-2.6K	Ohio MVECA FY26 schedule via extract; columns unlabelled; whether "per student" means enrolled or IEP students is unknown	Open (CT7)
PowerSchool Special Programs; Embrace	18,893/yr(<i>one/district</i>); 13K/yr (another)	Search summaries	Open (CT7)
Trinity (University Startups)	Not published	60-day money-back guarantee only	Unknown
IEP Compass (parent app)	\$9.99/mo or \$79.99/yr; 14-day trial	App Store	Confirmed
My IEP Hero (parent app)	\$9.99-29.99/mo	App Store / vendor (found by Val)	Confirmed
YouScience family test	~\$49	MO E4	Scout-cited
CollegeCountdown	Free tier; paid tiers not shown	Vendor	Vendor claim

Item	Price	Basis / caveat	Verdict
CollegeVine human mentoring	~\$1,300 on average	2026 third-party review	Confirmed
Independent educational consultant	"Just under \$140/hour" for many; multi-year package ~1/9 of \$40K ($\approx$ \$4,400, derived)	IECA FAQ, fetched by both comparator and validator	Confirmed
Consultant ranges	140 – 400/hr; 6,500 typical package	Published by interested sellers	Popularity tier
Pre-ETS national pot	~\$0.5-0.65B/yr	15% of a ~\$3.5-4.3B VR grant; GAO-18-502 FY2016 reserve of \$465M	Confirmed range
SchoolLinks revenue	~\$20-30M (2023 estimates)	Forbes \$30M; Latka \$23.3M; unaudited	Open, low stakes (CT9)

2c. Payers (for both specialists)

Payer	Evidence they pay for software in this space	Strength
District (counseling / CCR budget)	Yes: \$3-6 per student CCR licenses, multi-year, via RFP; Seattle's Naviance came from a tech levy (not re-verified)	High
State education agency	Yes: statewide CCR contracts in WI, FL, UT, WA; this pre-empts district-level sales	High
District special ed	Yes for IEP systems, TAGG (\$5 per set) and Trinity (paid; scale and price unverified). For student-facing survey tools the default is \$0	Medium; open (CT6)
State VR / Pre-ETS	Yes, but only career-exploration video (Virtual Job Shadow: NV statewide, co-purchased with the state education agency). No planner purchase found. Mostly buys human services. Anti-supplanting limit may block overlap with district CCR (untested). Some agencies are under budget stress (KY missed the 15% requirement in FFY2022 and closed all Order of Selection categories in May 2025)	Medium; see CT1, CT4
Pre-ETS providers	MA (~40 providers) and MN were funded to buy software	Medium (extracts)
Families	They pay for parent-side IEP apps (80 – 360/yr) and human counseling (140+/hr, concentrated in higher-income families seeking selective 4-year admission). No willingness-to-pay evidence for a <i>student-facing transition planner</i>	Open (CT5)
Colleges / recruiting marketplaces	Fund free student tools (CollegeVine, Naviance/Intersect); privacy-risky with disability data	High (structure)
Nonprofits	They <i>build</i> free tools (iTransition, CareerVillage); no evidence they buy licenses	Gap (GP7)
Medicaid waivers	No: they fund human services for a narrow I/DD population, not software (MO D5)	Finding (negative)
Student as payer; individual counselor or case manager	No evidence either way	Gap (GP13)

2d. Market-size figures (lift with the stated caveat)

Figure	Value	Caveat
IDEA students aged 3-21	7,525,941 (2022-23)	NCES Digest 204.50. Primary
IDEA students aged 14-21	~ 2.4M (2,414,034)	Leans toward this on arithmetic; the exact figure is a search extract (CT2). Do not cite NCLD's 4.1M or its 4.4% reach.
Pre-ETS-eligible (IEP + 504)	~3.1M	Hechinger 2025 (dataset not named)
Pre-ETS served	~180-300K (~6-10% reach)	Range only (CT3)
Recent graduates not at a 4-year college that fall	~58%	BLS 2024 cohort
District CCR total spend	"Tens to low hundreds of \$M"	Scout inference; no published figure (GP3)
Transition-specific spend	Unknown	Not in any public source (AG14, GP4)
11th-12th grade planning population (headcount)	Not reported by this pass	See gaps

3. Contradictions

From the validator (CONTESTED claims)

X1 (AG1). Is the CCR / IEP-transition seam "joined only by hand"?

- *Comparator*: no CCR platform offers an IEP transition module or a merged timeline. Counselors and special-ed staff work in silos, and procurement is siloed too (MS C3, Tyton 2026: 90% want cross-pillar partnerships, 14% pursue them). High that the seam is structural, medium that no vendor bridges it.
- *Validator*: the CCR-vendor part holds on independent checks. "No vendor bridges it" fails:
 - Trinity joins IEP transition to school and training recommendations from the special-ed side.
 - CollegeCountdown joins an IEP/504 hub to college search, aid and deadlines from the family side.
 - The Tyton report has no special-ed pillar and is vendor-sponsored (Renaissance), so it doesn't support a CCR/special-ed procurement silo.
- *Status*: **Broad version disproven. Narrow version stands.**
 - Products that bridge the seam in software do exist, per their own published materials.
 - Still standing: no *CCR platform* has an IEP module, and no product merges *dated IEP milestones* with aid and application deadlines.
 - How deep Trinity's and CollegeCountdown's features go is unverified (no demos).

X2 (AG8). Is per-school accommodations data "not embedded anywhere", and is CeDaR "the only national dataset"?

- *Comparator*: no CCR platform shows per-school accommodations. CeDaR is the only national database (~600 institutions with detail, degree-granting only).
- *Validator*:
 - CollegeCountdown advertises CeDaR-powered support tiers, and Trinity filters by disability support services.
 - The K&W Guide (Princeton Review, 350+ colleges with support-level codes) and Think College Search (300+ ID programs, with deadlines and aid eligibility) are also national datasets, though narrower in scope.
 - CeDaR's own page says "nearly 700", not ~600.
- *Status*: **"Only dataset" and "not embedded anywhere" are disproven** by the primary pages above.
- Still standing, confirmed:
 - No *district CCR platform* shows per-school accommodations.
 - CeDaR excludes trade and non-degree schools ("planned").
 - GAO-24-105614 finds that self-advocacy guidance doesn't reach transitioning students.
- The need is real. Per-school accommodations data *for trade and certificate schools* is the part nobody covers.

X3 (AG13). Is school counselor capacity thin for the target population?

- *Comparator*: the national ratio is 372-376:1 against ASCA's 250:1, and counselors spend a minority of their time on postsecondary work.
- *Validator*: the same ASCA release (Feb 2026) says the *high-school* ratio (195-224:1, from the 35 states reporting it) meets 250:1 for the first time. The shortfall is in elementary and middle schools (571-694:1).
- *Status*: **Unresolved as a claim about 11th-12th graders.**
 - Don't use 372:1 as evidence of thin high-school capacity.
 - The "minority of time on postsecondary" sub-point is unverified, and it is the part that could still make the case.

Carried from the comparator (validator CONFIRMED the comparator's treatment)

- **CT1. Has VR bought student-facing software? Partly resolved.** Yes for career-exploration video (NV, MA, MN; GA weak). Still open:
 - any VR purchase of a college, aid or transition *planner*;
 - dollar amounts;
 - whether Nevada renewed after 2022;
 - whether the anti-supplanting limit bites in practice.
- **CT2. IDEA students aged 14-21: ~2.4M vs ~4.1M. Leans toward 2.4M, not confirmed.**
 - NCLD's 4.1M would be ~55% of all IDEA students packed into 8 of 19 age-years, so it is implausible.
 - The implied reach (~7.6%) carries the same caveat.
- **CT3. Pre-ETS served: 295K (Hechinger 2023) vs 204K (RSA FFY21-22) vs 183K (NCLD RSA-911 PY2022). Open.**
 - Use the range.
 - NCLD's own caveat: RSA-911 may undercount.
- **CT4. Is Pre-ETS "often underspent"? Resolved in favour of "the story has faded".**

- Under-spending states fell from 23 (2021) to 9 (FY2022) (NCLD p.8, read directly).
- The ~\$196.5M FY25 relinquishment was returned for lack of state match. It is not unspent Pre-ETS money.
- FY25 reallocation requests (339.3M) *exceeded the funds available* (214.6M).
- Don't present idle Pre-ETS money as a funding source.
- The "\$146M FFY2026" figure is flagged; see section 5.
- **CT5. Will families pay? Open.**
 - *Parents-motivated side*: NLTS 2012, IN.
 - *Near-zero side*: MO G; CO 3.1.
 - The validator's new evidence undercuts "near-zero": two paid parent IEP apps, CollegeVine's ~\$1,300 mentoring, and freemium CollegeCountdown.
 - There is still no willingness-to-pay data point for a student-facing transition planner.
- **CT6. Does special ed pay for transition software? Open.**
 - "The default is \$0" holds for student-facing survey tools.
 - Trinity and TAGG show paid district special-ed purchases exist, but their scale and Trinity's price are unverified.
- **CT7. IEP-software cost per student: open.**
 - The per-student basis is unknown (enrolled vs IEP students).
 - Treat the ~\$35-160M IEP-software market estimate as unsupported.
- **CT8. IECA consultant rates: resolved** (see table 2b).
- **CT9. SchoolLinks revenue: open, low stakes** ("~\$20-30M, unverified, 2023").

Minor discrepancies (not material):

- Seattle has likely moved from Naviance to SchoolLinks.
- The iTransition partner-state lists are incomplete, not conflicting.
- Xello reach figures (~8M students vs 9M students and educators) are vendor self-reports on different bases.
- The date of Paper's MajorClarity acquisition is unsettled (2022 per Paper, Feb 2023 per EdWeek).

4. Gaps (still unknown after this pass)

1. **Expressed demand for a tool** from transition coordinators, special-ed teachers or parents. The need is documented; "I wish this existed" evidence is not. Forums weren't searchable (GP1).
2. **Procurement or RFP language** asking for an IEP-aware planner. None found (GP2).
3. **Total district CCR spend.** Only inference exists (GP3).
4. **Special-ed transition-specific spend.** None published; the national study reports in 2028 (GP4).
5. **Family willingness to pay** for a student-facing planner, family spend on private counseling, and rates for private transition consultants. The "26% use a private counselor" figure comes from an old survey of high achievers (GP5).
6. **A VR purchase of a planner, and RSA's view on whether one is allowable.** Neither exists in evidence (GP6).
7. **Nonprofits as buyers.** No evidence (GP7).
8. **Transition-coordinator headcount and caseloads.** No national data. This blocks sizing any per-seat practitioner model (GP8).

9. **The national Indicator 13 compliance average.** The ED file returned 403 (GP9).
10. **Competitor depth beyond public pages.** No demos. Plain-language and reading-level (e) couldn't be scored for any vendor. This now matters most for **CollegeCountdown and Trinity** (GP10).
11. **Adoption and usage** for iTransition, WiTransition, T-Folio, Trinity (only self-reported) and CollegeCountdown (unverified). It is also unknown whether the USC games are maintained or Nevada renewed (GP11).
12. **A current national survey of IEP families.** The latest is NLTS 2012 (GP12).
13. **The student as payer, and an individual counselor or case manager as buyer.** No evidence (GP13).
14. **Whether colleges use the Summary of Performance.** Minor (GP14).
15. **Single-state launch.** Not researched, by design. **Lead flag:** a product needs per-state aid and deadline data, and states pre-empt districts with statewide CCR contracts. A later pass will likely need one-state evidence (state CCR contract, state VR Pre-ETS delivery model, state promise aid). That needs an internal decision on which state, kept separate from the prototype student's identity (GP15).
16. **The 11th-12th grade headcount** (the general planning population the brief asked to size separately). This pass didn't report it. It is easy to fill from NCES enrollment tables in a later pass.
17. **Primary inputs seen only as search extracts:** the ED FY2026 Rehabilitation CJ (2,414,034; 120,119 referrals, 78% from schools), RSA's 203,638, the 47th IDEA Annual Report, and the Nevada, Utah, MN and Ohio records (GP16).

5. Flagged: insufficient sourcing

No whole claim came back UNSOURCED (validator tally: 0 of 24). The items below are sub-claims or figures inside CONFIRMED or CONTESTED claims. Either the validator couldn't find a source for them, or the comparator already marked them as inference or derived. They are listed so a deep pass knows what still needs real evidence. **Do not cite any of them as fact.**

#	Item	Why flagged	Origin
F1	"\$146M relinquished in FFY2026 by 15 agencies"	Not on NCLD p.7 or in Appendix C, where it was attributed. Location unknown	CT4 (Val)
F2	"Free transition tools cap price"	No source shows a buyer choosing free over paid. Trinity's paid sales point the other way	AG2 (Val)
F3	"State CCR deals up to ~\$9-10 per student"	One derived number from a Utah budget <i>request</i> , not a contract	AG9 (Val)
F4	Florida Xello ~\$2.8M/yr "from two separate sources"	Traces to The Drill Down alone; no FDOE or vendor figure	AG9 (Val)
F5	Xello "\$4.95 list"	Not re-found; the Park Hill quote shows \$5.45 regular	AG9 (Val)
F6	Tyton 2026 as evidence of a CCR/special-ed procurement silo	The report's pillars exclude special ed, and a vendor sponsored it	AG1 (Val)
F7	"Counselors spend a minority of time on postsecondary work"	Not verified	AG13 (Val)
F8	A planning app is Pre-ETS-allowable if positioned as a required activity	Scout inference in MO and IN; no RSA text	AG3, GP6

#	Item	Why flagged	Origin
F9	"Standalone CCR is hard to sustain"	Shared inference, not a sourced finding. The Cialfo restructuring and the College Conductor postmortem weren't re-checked	AG10
F10	IEP-software market ~\$35-160M/yr	Popularity-tier prices times scout arithmetic	CT7
F11	~\$4,400 consultant package	Arithmetic on the IECA wording; IECA doesn't print it	CT8
F12	NCLD/WestEd "37% of enrolled LD students attend vocational/technical schools"	Not verified by the validator	AG6
F13	NLTS parent-information items on training options and financial aid	Not re-found by the validator (the broader NLTS gap is confirmed)	AG7
F14	USC games "dormant"	Not verified	AG15
F15	Seattle funded Naviance from a tech levy, not special ed	Not re-verified	AG1
F16	CollegeCountdown feature list, "2,400+ reviews", "+127 SAT points"; Trinity's "more than a dozen states"	Vendor self-report only	Val
F17	2,414,034 (IDEA 14-21), RSA 203,638, ED CJ 120,119 referrals / 78%	Search extracts only; the exact numbers weren't found on a fetchable page	CT2, AG5, GP16
F18	NCLD 4.1M students aged 14-21 and 4.4% reach	Implausible on arithmetic; don't cite without the CT2 caveat	CT2
F19	"26% of families use a private counselor"	An old survey of high achievers only	GP5

6. Budget note

D-002 (`memory/decisions.md`) is still a **DRAFT pending Sharon's confirmation**:

- The original pilot ceiling was **≤150k tokens**.
- The first real pilot came in at 260,863, so the proposed revision is **≤300k for a 3-angle topic**, scaled proportionally by angle count.
- This run used **4 angles**, so the proportional draft ceiling is **≈400k** (the same figure the 2026-09-24 data point used).

The orchestrator logs the actual total in `budget-ledger.md` and should compare it against ≈400k (and note it against the original 150k draft).

Lead-stage variance: none. The lead did synthesis only (4 local reads, no web fetches).